

DANNY WASER

AI/ML ENGINEER | AI SYSTEMS, LLM & AUTOMATION

+41 (0) 78 911 98 34 | danny@waser.tech | Lausanne, CH

Availability: immediate

GitHub | LinkedIn | HF | waser.tech



PROFILE

AI/ML Engineer building AI systems end-to-end : from complex problems to production, covering data pipelines, LLM inference and fine-tuning, RAG, multi-agent orchestration, speech and intelligent automation. 57 merged PRs across 12+ open-source repositories (Coqui-STT, AutoAWQ, HuggingFace) ; 10% WER reduction on French speech models ; local inference of 70B+ models on a dual-GPU workstation. Direct C-suite collaboration : a functional prototype in under 2 weeks per feature. The red thread of my work : the interface between high-volume information and the people who act on it.

PROFESSIONAL COMPETENCIES

DATA SCIENCE & ANALYSIS:

- Prescriptive and predictive analytics : machine learning, predictive models, algorithms, decision support
- Data exploration and preparation : multi-source collection, preparation, truthfulness checking, visualization
- Data lifecycle : analysis tools, results interpretation, new data sources

DEVELOPMENT & ARCHITECTURE:

- Application design and development : development, debugging, documentation, integration, commissioning
- Architecture design : interoperability, scalability, security, vulnerability management
- Application design : data structures, modeling languages, iterative approach

MANAGEMENT & INNOVATION:

- Project and portfolio management : planning, resources and budget, cost-delay optimization, delivery
- Strategic business alignment : business needs, enterprise architecture, process efficiency, strategic decisions
- Innovation : creative solutions, new concepts, innovative mindset

TECHNICAL SKILLS

System Administration:

Data Center Operations, Headscale

Databases:

ETL, MariaDB, MongoDB, MySQL, PostgreSQL, SQL Server, SQLite, FTS5

Office & Management:

MS Office, Microsoft 365, LibreOffice

CMS:

Wordpress, Django CMS

Design Patterns:

MVC, MVVM

Framework:

Angular, Bootstrap, Django, jQuery, Node, PyTorch, React, Vue, TensorFlow, FastAPI, Hermes Agent

Machine Learning:

Deep learning, Machine learning, NLP, MLOps, AWQ Quantization, GPTQ Quantization, GGUF Quantization, PyTorch Internals, llama.cpp, vLLM, Multi-Agent Orchestration, RAG Pipelines, LoRA Fine-tuning, SFT Fine-tuning, Hugging Face Ecosystem, Prompt Engineering, DeepSpeech, Coqui STT, Whisper, OpenVoice, Voice Activity Detection, FFmpeg, OmniRoute Gateway, crawl4ai, Docling, Docling Serve, OpenCV, Pillow, pi, oh-my-pi, opencode, GitHub Copilot

OS:

Android, iOS, Linux, OSX, Unix, Windows, Windows serveur, systemd

Security:

Antivirus, Firewall, GPO, Spyware, SSL, VPN

Virtualisation:

HyperV, VirtualBox, VMWare, VMWare Workstation

Web Service:

REST

3D & Game Development:

Godot, Ursina

Processes:

BPM

Testing:

BDD, TDD

Server Applications:

Apache Tomcat, Apache HTTP Server, SearXNG, Freqtrade

BI & Data Science:

Tableau

Blockchain:

Ethereum, Solana

Cloud:

AWS, Azure, Firebase, Google Cloud Platform

Data Science:

Jupyter, Pandas, NumPy, Data Pipelines, SFT Dataset Curation

DevOps:

Ansible, CD, CI, Docker, GitHub Actions, GitLab, Grafana, Kubernetes, Prometheus, YAML, GNU, Arch Linux, Manjaro, NVIDIA GPU Server, CUDA Toolkit, Git & CI

Programming Interfaces (API):

DirectX, DOM, Vulkan, MCP, A2A

ITSM & Services:

JIRA

Programming Languages:

Bash, C++, CSS, HTML, Javascript, JSON, PHP, Python, SH, Shell, SQL, TypeScript, XML, CUDA C++, Triton Kernels, xonsh

Hardware & Equipment:

Desktop, Laptop, Imprimante, Peripherals, Routeurs, Smartphone, tablette, Switch

Networks, Protocols & Telecom:

3G, 4G, 5G, DHCP, Ethernet, HTTP, Internet, IP, LAN, WAN, POP3, SMTP, TCP, VoIP, Wifi, WebSocket

Testing & Quality:

Selenium, Playwright

Video Conferencing:

Teams, Zoom, kMeet

Web Scraping & Automation:

BeautifulSoup4, Scrapy

Governance:

Lean, RGPD

Projects:

Agile, DevOps, Kanban, Scrum

CERTIFICATIONS

Google Advanced Data Analytics Professional Certificate (Google @ Coursera | 10.2023 to 11.2023)

LANGUAGES

French (C2) | English (C1) | Spanish (C1) | German (B1)

PROFESSIONAL EXPERIENCE

Idéal Isolation Sàrl & DF Innovation Sàrl

09.2024 to 03.2026

Data Scientist

- Operational analytics : bottleneck analysis of quotation/invoicing workflows; data pipelines + predictive models for resource allocation.
- Generative AI R&D : custom RAG pipelines, fine-tuned LLMs for document understanding and synthesis.
- Algorithmic trading : ML signal extraction from market data, backtesting infrastructure, performance analysis, production bot (live + dry-run).
- Blockchain : Solana smart-contract prototypes for automated market operations.

Synthflow AI

01.2024 to 03.2024

Machine Learning Engineer

- Multi-agent conversational systems : production voice-call systems with specialized LLM agents in tandem for telephony.
- Real-time dialogue management : LLM-based dialogue managers, dynamic prompt engineering, sub-second response in live calls.
- Rapid R&D cycle : concept to functional prototype in <2 weeks per feature; direct C-suite collaboration.

Mozilla & Coqui-AI

11.2021 to 12.2022

Speech recognition developer

- 12 merged PRs in coqui-ai/STT : memory management in batch training [#2205](#), TF graph export [#2222](#), TFLite export [#2221](#), LM optimization [#2116](#).
- CommonVoice FR maintenance ([PR #168](#)) : DeepSpeech training image on STT framework.
- Published French speech models : 2x higher-quality augmented data, -10% average Word Error Rate.
- DeepSpeech → STT migration : led the French training pipeline migration.

Waser Technologies

01.2020 to 07.2023

AI Developer & Founder (Entrepreneurial)

- CUDA kernel optimization : 2 merged PRs in [AutoAWQ #573](#) & [AutoAWQ_kernels #29](#); quantization build pipeline.
- Multi-modal translator ([wasertech/Translator](#)) : batch CLI, NLLB-200, 4 merged PRs.
- Real-time speech pipeline ([waser-technologies/listen](#)) : Faster-Whisper + VAD, low-latency transcription, OpenVoice v2.
- Multi-agent LLM orchestrator ([waser-technologies/assistant](#)) : production assistant on local hardware (96GB RAM, 2xRTX Titan); natural-language ops, on-premise inference.
- Production MLOps : containerized deployments with GPU passthrough (Docker Compose), service supervision (systemd, s6-overlay), model lifecycle and performance monitoring on local inference stack.
- HuggingFace transformers ([PR #42864](#)) : merged activation-init fix.

EDUCATION

Commercial Diploma (CFC + MPC)

École de Culture Générale et de Commerce | 08.2014 to 07.2017

PORTFOLIO

Hugging Face Profile | Official Hugging Face profile hosting open-source LLMs, speech models, and specialized ML datasets.

Assistant | Multi-agent LLM shell orchestrating specialised agents for system ops, web research, and file editing via a natural language prompt.

Listen | Real-time speech-to-text system with server-client architecture using Faster-Whisper and VAD for efficient audio transcription.

Say | Text-to-speech tool generating high-quality customizable voices from 10-second audio samples using OpenVoice v2.

Translator | Batch-friendly translation CLI leveraging Hugging Face translation pipelines (NLLB-200 and compatible models) with full support for metadata-preserving PO file handling.

OPEN SOURCE CONTRIBUTIONS

- **Mozilla Common Voice**
French STT Model Training Pipeline
- **huggingface/transformers** | 1 PR (164,514 stars)
fix: Initialize ApertusMLP's xielu activation using `torch\_dtype` (+7/-2)
- **django-cms/django-cms** | 1 PR (10,665 stars)
fix: typos in French translations (+1/-1)
- **coqui-ai/STT** | 12 PRs (2,606 stars)
Bring back importer for TrainingSpeech (+293/-0)
Add excluded sentences to LM (+119/-17)
Fix CCPMF importer (+97/-31)
Make lm_optimizer.py as python module (+71/-47)
Workaround batch size memory test (+61/-7)
export to tflite using python37 venv (+32/-13)
Fix LM optimization script importing non existing function and add CI (+37/-1)
Fix importers save_excluded_max_sec_to_disk flag (+31/-4)
Update BUILDING.rst (+8/-8)
Updated missing checkpoint path error message (+11/-1)
Remove `convert\_graphdef\_memmapped\_format` (+0/-4)
[FIX]: lm_optimize module import from training utils (+1/-1)
- **casper-hansen/AutoAWQ** | 1 PR (2,349 stars)
Let installed PyTorch decide required version number (+2/-1)
- **ixray-team/ixray-1.6-stcop** | 3 PRs (442 stars)
Support common cross-platform rendering pipeline thru Vulkan (+3186/-42)
Fix language reload and translate inventory items when switching langu (+173/-12)
Add support for preferred fallback language (+90/-4)
- **coqui-ai/STT-models** | 6 PRs (153 stars)
Update cv-fr v0.8 and add v0.9 (+620/-4)
New release 0.8.3 for French STT (+600/-0)
Fix broken link for chatio (+1/-1)
Update MODEL_CARD.md (+1/-1)
Update MODEL_CARD.md (+1/-1)
Update MODEL_CARD.md (+1/-1)
- **coqui-ai/STT-examples** | 3 PRs (132 stars)
Add Django streaming example (+1492/-0)
Update README.md (+3/-3)
Update README.rst (+1/-0)
- **casper-hansen/AutoAWQ_kernels** | 1 PR (80 stars)
Update build process for flexibility and torch version compatibility (+26/-9)
- **EsupPortal/Esup-Pod** | 1 PR (48 stars)
Use -STT (+21/-21)
- **coqui-ai/stt-model-manager** | 4 PRs (26 stars)
Fix security vulnerabilities (+3898/-2587)
Version 0.0.21-GM (+88/-26)
Loosen STT version requirement (+4/-4)
Loosen `coqpit` (+1/-1)
- **dperelman/gamepad-util** | 1 PR (15 stars)
Update identify_evdev.py (+1/-1)
- **ixray-community/ixray-unofficial-patch** | 1 PR (2 stars)
Update README with mod loading instructions (+3/-0)

PERSONAL INTERESTS

Fan of astronomy, space opera, mountain trails, conscious slam, epic adventures