

DANNY WASER

MACHINE LEARNING ENGINEER | DATA SCIENTIST

+41 (0) 78 911 98 34 | danny@waser.tech | Lausanne, CH

Disponibilité : immédiate

GitHub | LinkedIn | HF | waser.tech



PROFILE

Machine Learning Engineer & Data Scientist: end-to-end ML solutions — from data exploration and analysis (e-CF D.7) to production deployment (B.4) — including data pipelines, RAG, LLM fine-tuning, speech and automation. 57 merged PRs across 12+ open-source repositories (Coqui-STT, AutoAWQ, HuggingFace); 10% WER reduction on French speech models; local inference of 70B+ models on a dual-GPU workstation. Direct C-suite collaboration: functional prototype in under 2 weeks per feature.

PROFESSIONAL COMPETENCIES

PLANIFIER: A.5 Conception de l'architecture • A.7 Veille technologique • A.1 Systèmes d'information et alignement stratégique métier • A.3 Mise en place d'un plan d'activités • A.4 Planification des produits/services

DÉVELOPPER: B.6 Ingénierie système • B.1 Conception et développement d'applications • B.2 Intégration des composants • B.3 Tests • B.4 Déploiement de la solution

EXPLOITER: C.2 Support aux changements • C.4 Gestion des problèmes • C.5 Gestion des systèmes • C.1 Support utilisateur • C.3 Fourniture de services

FACILITER: D.7 Science des données et analyse • D.1 Développement d'une stratégie de sécurité de l'information • D.2 Développement d'une stratégie pour la qualité des TIC • D.10 Gestion de l'information et de la connaissance • D.4 Achats

GÉRER: E.1 Développement prévisionnel • E.2 Gestion des projets et du portefeuille de projets • E.5 Amélioration des processus • E.6 Gestion de la qualité des TIC • E.9 Gouvernance des systèmes d'information

TECHNICAL SKILLS

Administration systèmes:

Data Center Operations

Base de données:

ETL, MariaDB, MongoDB, MySQL, PostgreSQL, SQL Server, SQLite, FTS5

Bureautique / Gestion:

MS Office, Microsoft 365

CMS:

Wordpress, Django CMS

Design Patterns:

MVC, MVVM

Framework:

Angular, Bootstrap, Django, jQuery, Node, PyTorch, React, Vue, TensorFlow, FastAPI, Hermes Agent

Machine Learning:

Deep learning, Machine learning, NLP, AWQ Quantization, GPTQ Quantization, GGUF Quantization, PyTorch Internals, llama.cpp, vLLM, Multi-Agent Orchestration, RAG Pipelines, LoRA Fine-tuning, SFT Fine-tuning, Hugging Face Ecosystem, Prompt Engineering, DeepSpeech, Coqui STT, Whisper, OpenVoice, Voice Activity Detection, FFmpeg, OmniRoute Gateway, crawl4ai, Docling, Docling Serve, OpenCV, Pillow, pi, oh-my-pi, opencode, GitHub Copilot

OS:

Android, iOS, Linux, OSX, Unix, Windows, Windows serveur, systemd

Sécurité:

Antivirus, Firewall, GPO, Spyware, SSL, VPN

Virtualisation:

HyperV, VirtualBox, VMWare, VMWare Workstation

Web Service:

REST

Processus:

BPM

Test:

BDD, TDD

Applications serveur:

Apache Tomcat, Apache HTTP Server, SearXNG, Freqtrade

BI & Data science:

Tableau

Blockchain:

Ethereum, Solana

Cloud:

AWS, Azure, Firebase, Google Cloud Platform

Data Science:

Jupyter, Pandas, NumPy, Data Pipelines, SFT Dataset Curation

DevOps:

Ansible, CD, CI, Docker, GitHub Actions, GitLab, Grafana, Kubernetes, Prometheus, YAML, GNU, Arch Linux, Manjaro, NVIDIA GPU Server, CUDA Toolkit, Git & CI

Interface de programmation (API):

DirectX, DOM

ITSM & Services:

JIRA

Langages informatiques:

Bash, C++, CSS, HTML, Javascript, JSON, PHP, Python, SH, Shell, SQL, TypeScript, XML, CUDA C++, Triton Kernels, xonsh

Matériel & Equipement:

Desktop, Laptop, Imprimante, Périphériques, Routeurs, Smartphone, tablette, Switch

Réseaux, protocoles & Télécommunications:

3G, 4G, 5G, DHCP, Ethernet, HTTP, Internet, IP, LAN, WAN, POP3, SMTP, TCP, VoIP, Wifi, WebSocket

Test & Qualité:

Selenium

Visioconférence:

Teams, Zoom

Gouvernance:

Lean, RGPD

Projet:

Agile, DevOps, Kanban, Scrum

CERTIFICATIONS

Google Advanced Data Analytics Professional Certificate (Google @ Coursera | 10.2023 - 11.2023)

LANGUAGES

French (C2) | English (C1) | Spanish (C1) | German (B1)

PROFESSIONAL EXPERIENCE

Idéal Isolation Sàrl & DF Innovation Sàrl

09.2024 - 03.2026

Data Scientist

- Operational Analytics (e-CF D.7): Strategic analysis of bottlenecks in quotation and invoicing workflows; developed data pipelines and predictive models to optimize resource allocation and reduce processing delays — in direct collaboration with management on business requirements.
- Generative AI R&D (e-CF B.1): Designed and developed custom RAG pipelines and fine-tuned LLMs for domain-specific document understanding and synthesis.
- Algorithmic Trading: Developed quantitative trading strategies using ML-driven signal extraction from market data, with backtesting infrastructure and performance analysis.
- Blockchain: Prototyped Solana smart-contracts for automated market operations.

Synthflow AI

01.2024 - 03.2024

Machine Learning Engineer

- Multi-agent conversational systems (e-CF B.1): Designed and prototyped production voice-call systems deploying multiple specialized LLM agents in tandem for telephony.
- Real-time dialogue management: Iterative implementation of LLM-based dialogue managers with dynamic prompt engineering for sub-second response times in live calls.
- Rapid R&D cycle (e-CF E.2): From concept to functional prototype in under 2 weeks per feature, working directly with C-suite on product requirements.

Mozilla & Coqui-AI

11.2021 - 12.2022

Speech recognition developer

- 12 merged PRs in coqui-ai/STT across 1.5 years: Fixes to memory management during batch training [#2205](#), TF graph export optimization [#2222](#) ; removed memmapped format, TFLite export pipeline [#2221](#), LM optimization module refactoring [#2116](#), [#2203](#), importer fixes [#2148](#), [#2111](#), and CI improvements [#2285](#).
- CommonVoice FR maintenance ([PR #168](#)): Reimplemented the DeepSpeech model training image using the STT framework to continue French transcription capacity after DeepSpeech deprecation.
- Published French speech models (github.com/wasertech/commonvoice-fr/releases): Trained new French models using 2x higher-quality augmented data, reducing average Word Error Rate by 10%.
- DeepSpeech → STT migration: Led the migration of the French-language training pipeline from deprecated DeepSpeech to the STT codebase.

Waser Technologies

01.2020 - 07.2023

AI Developer & Founder (Entrepreneurial)

- CUDA kernel optimization: Merged PRs in [AutoAWQ #573](#) and [AutoAWQ_kernels #29](#), updated build process for Torch version compatibility and flexible quantization kernel compilation.
- Multi-modal translator (github.com/wasertech/Translator): Batch translation CLI with PO file support, multi-language detection, and NLLB-200 integration. 4 merged PRs covering locale handling, type safety, and PO workflow.
- Real-time speech transcription pipeline (gitlab.com/waser-technologies/technologies/listen): End-to-end client-server architecture combining Faster-Whisper with VAD for latency-optimized live transcription; paired with OpenVoice v2 for few-shot voice cloning and TTS.
- Multi-agent LLM orchestrator (gitlab.com/waser-technologies/technologies/assistant): production assistant running on local hardware (96GB RAM, dual RTX Titan) orchestrating specialized agents for system ops, web research, and file editing via natural language, with on-premise inference serving.
- HuggingFace transformers ([PR #42864](#)): Merged fix for ApertusMLP xielu activation initialization using torch_dtype.

EDUCATION

Commercial Diploma (CFC + MPC)

École de Culture Générale et de Commerce | 08.2014 - 07.2017

PORTFOLIO

Hugging Face Profile | Official Hugging Face profile hosting open-source LLMs, speech models, and specialized ML datasets.

Assistant | Multi-agent LLM shell orchestrating specialised agents for system ops, web research, and file editing via a natural language prompt.

Listen | Real-time speech-to-text system with server-client architecture using Faster-Whisper and VAD for efficient audio transcription.

Say | Text-to-speech tool generating high-quality customizable voices from 10-second audio samples using OpenVoice v2.

Translator | Batch-friendly translation CLI leveraging Hugging Face translation pipelines (NLLB-200 and compatible models) with full support for metadata-preserving PO file handling.

OPEN SOURCE CONTRIBUTIONS

- **Mozilla Common Voice**
French STT Model Training Pipeline
- **huggingface/transformers** | 1 PR (164,276 stars)
fix: Initialize ApertusMLP's xielu activation using `torch\_dtype` (+7/-2)
- **django-cms/django-cms** | 1 PR (10,661 stars)
fix: typos in French translations (+1/-1)
- **coqui-ai/STT** | 12 PRs (2,604 stars)
Bring back importer for TrainingSpeech (+293/-0)
Add excluded sentences to LM (+119/-17)
Fix CCPMF importer (+97/-31)
Make lm_optimizer.py as python module (+71/-47)
Workaround batch size memory test (+61/-7)
export to tflite using python37 venv (+32/-13)
Fix LM optimization script importing non existing function and add CI (+37/-1)
Fix importers save_excluded_max_sec_to_disk flag (+31/-4)
Update BUILDING.rst (+8/-8)
Updated missing checkpoint path error message (+11/-1)
Remove `convert\_graphdef\_memmapped\_format` (+0/-4)
[FIX]: lm_optimize module import from training utils (+1/-1)
- **casper-hansen/AutoAWQ** | 1 PR (2,349 stars)
Let installed PyTorch decide required version number (+2/-1)
- **ixray-team/ixray-1.6-stcop** | 3 PRs (440 stars)
Support common cross-platform rendering pipeline thru Vulkan (+3186/-42)
Fix language reload and translate inventory items when switching langu (+173/-12)
Add support for preferred fallback language (+90/-4)
- **coqui-ai/STT-models** | 6 PRs (153 stars)
Update cv-fr v0.8 and add v0.9 (+620/-4)
New release 0.8.3 for French STT (+600/-0)
Fix broken link for chatio (+1/-1)
Update MODEL_CARD.md (+1/-1)
Update MODEL_CARD.md (+1/-1)
Update MODEL_CARD.md (+1/-1)
- **coqui-ai/STT-examples** | 3 PRs (132 stars)
Add Django streaming example (+1492/-0)
Update README.md (+3/-3)
Update README.rst (+1/-0)
- **casper-hansen/AutoAWQ_kernels** | 1 PR (80 stars)
Update build process for flexibility and torch version compatibility (+26/-9)
- **EsupPortal/Esup-Pod** | 1 PR (48 stars)
Use -STT (+21/-21)
- **coqui-ai/stt-model-manager** | 4 PRs (26 stars)
Fix security vulnerabilities (+3898/-2587)
Version 0.0.21-GM (+88/-26)
Loosen STT version requirement (+4/-4)
Loosen `coqpit` (+1/-1)
- **dperelman/gamepad-util** | 1 PR (15 stars)
Update identify_evdev.py (+1/-1)
- **ixray-community/ixray-unofficial-patch** | 1 PR (2 stars)
Update README with mod loading instructions (+3/-0)

PERSONAL INTERESTS

Fan of astronomy, space opera, mountain trails, conscious slam, epic adventures